

# Student Task Breakdown — 2-Week Sprint

## ⓘ Scope

**Duration:** Feb 17 – Feb 28 (2 weeks)

**Budget:** 8 hrs/week per student = 16 hrs each, 48 hrs total

**Goal:** Working backend API + clickable Figma prototype. By the end of this sprint, the frontend and backend can start talking to each other.

## ⌚ Supervision cadence

- **Monday AM:** Send task briefs (copy from this doc)
- **Wednesday:** Async check-in — each student posts a screenshot or code snippet of where they are
- **Friday:** Deliverable review — 15 min call or async demo

## Mario & Steven — Backend (Python)

### Week 1: Foundation (Feb 17–21)

**Goal:** Get TransformerLens running, understand the model, build the first API endpoint.

#### Task 1.1: Environment Setup (~2 hrs)

**Who:** Steven, Mario

**What to do:**

- Create a project folder `interpretative-interfaces-backend/`
- Set up a Python virtual environment (`python -m venv venv`)
- Install dependencies: `pip install transformer_lens flask numpy scikit-learn`
- Create `requirements.txt` with pinned versions
- Load GPT-2 small and run one forward pass to confirm it works:

```
from transformer_lens import HookedTransformer
model = HookedTransformer.from_pretrained("gpt2-small")
logits, cache = model.run_with_cache("Hello world")
print(logits.shape) # should be (1, 2, 50257)
print(cache["resid_post", 0].shape) # should be (1, 2, 768)
```

**Done when:** you can run the above snippet on your laptops without errors.

### Resources:

- TransformerLens install: <https://github.com/TransformerLensOrg/TransformerLens>
  - Python venv docs: <https://docs.python.org/3/library/venv.html>
- 

## Task 1.2: Work Through the "50 Lines" Tutorial (~3 hrs)

**Who:** Steven, Mario

### What to do:

- Open the tutorial notebook in Google Colab (free, no GPU needed for GPT-2 small)
- Run every cell. Don't just read ; type the code, look at the outputs
- For each section, write a 1-sentence comment explaining what it does
- **Key things to understand:**
  - `model.run_with_cache()` → how to get all activations
  - `cache["resid_post", layer]` → residual stream at a layer (token embeddings)
  - `cache["blocks.X.attn.hook_pattern"]` → attention weights
  - `model.to_str_tokens()` → tokenization
- Save your annotated notebook as `notebooks/tutorial-walkthrough.ipynb`

**Done when:** When you can explain in your own words: (1) what the residual stream is, (2) how to get a token's embedding at layer 5, (3) what attention patterns show.

### Resources:

- Tutorial: [How-to Transformer Mechanistic Interpretability in 50 Lines](#)
  - TransformerLens demo notebook:  
[https://transformerlensorg.github.io/TransformerLens/content/getting\\_started\\_mech\\_interp.html](https://transformerlensorg.github.io/TransformerLens/content/getting_started_mech_interp.html)
-

## Task 1.3: Tokenization Function + Notebook (1.5 hrs)

**Who:** Steven

**What to do:**

- Create notebooks/tokenization.ipynb
- Write a function that takes a string and returns a list of token objects:

```
def tokenize(text: str) -> list[dict]:  
    """Returns list of {index, token_str, token_id} for each token."""  
    tokens = model.to_tokens(text)           # tensor of token IDs  
    str_tokens = model.to_str_tokens(text)   # list of strings  
    return [  
        {"index": i, "token_str": s, "token_id": int(tokens[0, i])}  
        for i, s in enumerate(str_tokens)  
    ]
```

- Test with 3 different inputs (short sentence, longer sentence, sentence with unusual words)
- Note any surprises (subword tokens, spacing, special tokens like BOS)

**Done when:** Function works, tested with 3 inputs, outputs saved in the notebook.

---

## Task 1.4: Embedding Extraction Across All Layers (~2 hrs)

**Who:** Mario

**What to do:**

- Create notebooks/embedding\_extraction.ipynb
- Write a function that takes text and a token index, and returns that token's embedding at every layer:

```
def get_token_trajectory(text: str, token_index: int) -> list[dict]:  
    """Returns the residual stream vector for one token at each layer."""  
    _, cache = model.run_with_cache(text)  
    n_layers = model.cfg.n_layers  # 12 for GPT-2 small  
    trajectory = []  
    for layer in range(n_layers):  
        embedding = cache["resid_post", layer][0, token_index, :]  # shape  
        (768, )  
        trajectory.append({
```

```

        "layer": layer,
        "embedding": embedding.tolist() # convert tensor to list
    })
return trajectory

```

- Test: extract trajectory for the word "cat" in "The cat sat on the mat"
- Print the shape and first 5 values at each layer to verify they change

**Done when:** Function works, trajectory returned for all 12 layers, values visibly differ across layers.

---

## Task 1.5: Dimensionality Reduction (~2.5 hrs)

**Who:** Steven

**What to do:**

- Create notebooks/dimensionality\_reduction.ipynb
- Write a function that takes multiple token trajectories and reduces them to 2D using PCA:

```

from sklearn.decomposition import PCA
import numpy as np

def reduce_trajectories(trajectories: dict[int, list]) -> dict[int,
list[dict]]:
    """
    Input: {token_index: [768-dim embedding per layer]}
    Output: {token_index: [{layer, x, y} per layer]}

    All tokens share the same PCA space so trajectories are comparable.
    """

    # Collect ALL embeddings into one matrix for joint PCA
    all_embeddings = []
    labels = [] # (token_index, layer) pairs
    for token_idx, layers in trajectories.items():
        for layer_data in layers:
            all_embeddings.append(layer_data["embedding"])
            labels.append((token_idx, layer_data["layer"]))

    all_embeddings = np.array(all_embeddings)
    pca = PCA(n_components=2)
    coords_2d = pca.fit_transform(all_embeddings)

```

```

# Reconstruct per-token trajectories with 2D coordinates
result = {}
for i, (token_idx, layer) in enumerate(labels):
    if token_idx not in result:
        result[token_idx] = []
    result[token_idx].append({
        "layer": layer,
        "x": float(coords_2d[i, 0]),
        "y": float(coords_2d[i, 1])
    })
return result

```

- Test: reduce 2-3 tokens from the same sentence, plot the trajectories with matplotlib
- **Bonus:** Try UMAP (`pip install umap-learn`) as an alternative and compare

**Done when:** 2D scatter plot showing 2-3 token trajectories through 12 layers, each trajectory is a connected path of dots. Save the plot in the notebook.

## Resources:

- scikit-learn PCA: <https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>
  - UMAP docs: <https://umap-learn.readthedocs.io/en/latest/>
- 

## Task 1.6: Flask App Skeleton + /tokenize Endpoint (~2 hrs)

**Who:** Mario

**What to do:**

- Create `app.py` with a basic Flask server
- Move the tokenization function from the notebook into `model_utils.py`
- Load the model once at startup (global variable — it's fine for a prototype)
- Implement the first endpoint:

```

POST /tokenize
Body: {"text": "The cat sat on the mat"}
Response: {
    "tokens": [
        {"index": 0, "token_str": "The", "token_id": 464},
        {"index": 1, "token_str": " cat", "token_id": 3797},
        ...
    ]
}

```

```
]  
}
```

- Add CORS headers (`pip install flask-cors`) so a browser frontend can call it
- Test with `curl`:

```
curl -X POST http://localhost:5000/tokenize \  
-H "Content-Type: application/json" \  
-d '{"text": "The cat sat on the mat"}'
```

**Done when:** `curl` returns correct JSON. Server starts without errors.

### Resources:

- Flask quickstart: <https://flask.palletsprojects.com/en/stable/quickstart/>
  - flask-cors: <https://flask-cors.readthedocs.io/en/latest/>
- 

## Week 1 File Structure

```
interpretative-interfaces-backend/  
|__ venv/  
|__ requirements.txt  
|__ app.py                      # Flask server  
|__ model_utils.py               # tokenize(), get_token_trajectory(),  
reduce_trajectories()  
|__ notebooks/  
|   |__ tutorial-walkthrough.ipynb # Task 1.2  
|   |__ tokenization.ipynb        # Task 1.3  
|   |__ embedding_extraction.ipynb # Task 1.4  
|   |__ dimensionality_reduction.ipynb # Task 1.5  
|__ README.md                    # Setup instructions
```

---

## Week 2: API Completion (Feb 24–28)

**Goal:** All API endpoints working, example data generated for Polly, documentation complete.

---

### Task 2.1: /trace Endpoint (~3 hrs)

## Who: Steven

### What to do:

- Combine the embedding extraction + PCA functions into a single endpoint:

```
POST /trace
Body: {
    "text": "The cat sat on the mat",
    "token_indices": [1, 4]
}
Response: {
    "tokens": [
        {"index": 1, "token_str": " cat"},
        {"index": 4, "token_str": " the"}
    ],
    "trajectories": {
        "1": [
            {"layer": 0, "x": -2.34, "y": 1.56},
            {"layer": 1, "x": -1.89, "y": 2.01},
            ...
        ],
        "4": [
            {"layer": 0, "x": 0.45, "y": -0.78},
            ...
        ]
    },
    "pca_explained_variance": [0.34, 0.21]
}
```

- Include the PCA explained variance ratio so the frontend knows how much information the 2D view captures
- Test with `curl` and verify the coordinates change across layers

**Done when:** Endpoint returns well-structured JSON with 2D coordinates for all 12 layers.

Tested with 2+ different inputs.

---

## Task 2.2: /attention Endpoint (~3 hrs)

### Who: Mario

### What to do:

- Create an endpoint that returns attention patterns for a given layer and head:

```
POST /attention
Body: {
  "text": "The cat sat on the mat",
  "layer": 5,
  "head": 3
}
Response: {
  "tokens": ["The", " cat", " sat", " on", " the", " mat"],
  "attention_matrix": [
    [0.85, 0.05, 0.03, 0.02, 0.03, 0.02],
    [0.15, 0.45, 0.20, 0.08, 0.07, 0.05],
    ...
  ]
}
```

- The attention matrix is a square matrix (`n_tokens x n_tokens`) where `[i][j]` is how much token `i` attends to token `j`
- Extract from cache: `cache["blocks.{layer}.attn.hook_pattern"] [0, head]`
- **Bonus:** Add an optional parameter to return all heads for a layer at once

**Done when:** Returns correct attention matrix. Verified by checking rows sum to ~1.0 (they're softmaxed probabilities).

---

## Task 2.3: /predict Endpoint — Logit Lens (~3 hrs)

**Who:** Steven

**What to do:**

- Create an endpoint that shows what the model would predict at each layer (logit lens technique):

```
POST /predict
Body: {
  "text": "The cat sat on the mat",
  "token_index": 5
}
Response: {
  "predictions_by_layer": [
    {
      "layer": 0,
      "prediction": "The"
    },
    ...
  ]
}
```

```

    "layer": 0,
    "top_tokens": [
        {"token": " the", "probability": 0.08},
        {"token": " a", "probability": 0.05},
        {"token": ",", "probability": 0.04},
        {"token": " and", "probability": 0.03},
        {"token": " of", "probability": 0.02}
    ]
},
{
    "layer": 1,
    "top_tokens": [...]
},
...
]
}

```

- For each layer, project the residual stream through the unembedding matrix:

```

residual = cache["resid_post", layer][0, token_index] # (768,)
logits =
model.unembed(model.ln_final(residual.unsqueeze(0).unsqueeze(0))) # project to vocab
probs = torch.softmax(logits[0, 0], dim=-1)
top5 = torch.topk(probs, 5)

```

- Return top 5 predicted tokens and their probabilities at each layer

**Done when:** Can see predictions evolve across layers. Early layers should be vague/generic, later layers should converge on the actual next token.

## Resources:

- Logit lens explanation:  
<https://www.lesswrong.com/posts/AcKRB8wDpdaN6v6ru/interpreting-gpt-the-logit-lens>

## Task 2.4: Generate Example Data for Frontend (~2 hrs)

**Who:** Mario

**What to do:**

- Create examples/ folder
- Run all 4 endpoints with 3 different input texts and save the JSON responses:

1. "The cat sat on the mat" (simple, familiar)
  2. "Her name was Alex Hart. Tomorrow Alex" (the induction example from the tutorial)
  3. "The Eiffel Tower is located in the city of" (knowledge recall)
- Save as:
    - examples/example1\_tokenize.json
    - examples/example1\_trace.json
    - examples/example1\_attention.json
    - examples/example1\_predict.json
    - (repeat for examples 2 and 3)

**Done when:** 12 JSON files in examples/. Polly can use these as real mock data for the frontend without needing to run the backend.

---

## Task 2.5: API Documentation + README (~2 hrs)

**Who:** Steven, Mario

**What to do:**

- Update README.md with:
  - Setup instructions (clone, create venv, install deps, run)
  - All 4 endpoints documented (URL, method, request body, response format)
  - Example curl commands for each endpoint
- Create API.md with the full JSON schemas

**Done when:** Someone who wasn't on the project can follow the README, start the server, and successfully call all 4 endpoints.

---

## Week 2 File Structure (Final)

```
interpretative-interfaces-backend/
├── venv/
├── requirements.txt
├── app.py                                # Flask server with all endpoints
├── model_utils.py                         # All model functions
└── notebooks/
    ├── tutorial-walkthrough.ipynb
    └── tokenization.ipynb
```

```
|   └── embedding_extraction.ipynb  
|   └── dimensionality_reduction.ipynb  
└── examples/  
    ├── example1_tokenize.json  
    ├── example1_trace.json  
    ├── example1_attention.json  
    ├── example1_predict.json  
    ├── example2_*.json  
    └── example3_*.json  
└── API.md                                # Endpoint documentation  
└── README.md                             # Setup + usage instructions
```

---

## Polly — Frontend/Design (Figma)

### Week 1: Research + Wireframes (Feb 17–21)

---

#### Task P1.1: Study Existing Projects (~2 hrs)

##### What to do:

- Spend time with these 3 projects, take notes on what works and what doesn't:
  1. **Transformer Explainer** — <https://poloclub.github.io/transformer-explainer/>
    - Play with it for 15 min. Note: the Sankey flow diagram, the hover interactions, the temperature slider
  2. **LLM Visualization (Brendan Bycroft)** — <https://bbycroft.net/llm>
    - Play with it for 15 min. Note: the 3D animation, the sense of "flowing through" the architecture
  3. **Ecco documentation** — <https://ecco.readthedocs.io/en/latest/>
    - Look at the gallery/examples. Note: how token trajectories are visualized (2D scatter with connected paths)
- Write a short doc (1 page) with:
  - 3 things that work well across these projects
  - 3 things that are missing or could be better
  - 1 initial idea for how our project could be different

**Done when:** 1-page doc shared with the team.

##### Resources:

- Our Figma board with the examples of existing projects and brainstorming session recap.
- 

## Task P1.2: User Flow Wireframes (~3 hrs)

### What to do:

- In Figma, create low-fidelity wireframes (gray boxes, no styling) for the core user flow:
  1. **Start screen:** Text input field, "Analyze" button
  2. **Token selection:** Text is displayed as individual tokens. User clicks to select 1-3 tokens (they highlight)
  3. **Trajectory view:** 2D scatter plot showing selected tokens' paths through 12 layers. Each token is a different color. Layers are connected by lines. User can hover a point to see layer number and nearby words
  4. **Attention view:** Select a layer → see the attention heatmap (which tokens attend to which). Matrix or arc diagram
  5. **Prediction view:** Select a token position → see top-5 predicted words at each layer (the logit lens). Shows how predictions evolve from vague to specific
- Don't worry about visual polish yet — focus on layout and flow

**Done when:** 5 wireframe screens in Figma, connected with basic click-through navigation.

### Resources:

- API response schemas (coordinate with Mario/Steven for JSON structure)
  - Transformer Explainer for layout inspiration
- 

## Task P1.3: Interaction Sketches (~3 hrs)

### What to do:

- For the **trajectory view** (the most important screen), sketch 3 different representation ideas:
  1. **Scatter plot** — tokens as colored dots in 2D PCA space, lines connecting same token across layers (like Ecco)
  2. **Layer stack** — layers as horizontal rows, tokens move left/right as their position in semantic space shifts
  3. **One wild card** — your choice. Could be: circular/radial, 3D depth, organic/flowing, or something entirely new

- For each, annotate: how does the user select a layer? How do they compare tokens? Where do annotations go?

**Done when:** 3 sketches in Figma with annotations.

---

## Week 2: Clickable Prototype (Feb 24–28)

---

### Task P2.1: Build Clickable Prototype (~4 hrs)

**What to do:**

- Pick the strongest wireframe from Week 1 (or combine elements)
- Build a mid-fidelity Figma prototype with:
  - Real-ish data (use the JSON examples from Mario/Steven's examples/ folder)
  - Click-through flow from text input → token selection → trajectory view → attention view
  - Hover states on tokens and trajectory points
  - Color system: 2-3 colors for selected tokens, neutral for background
- Use Figma's native prototyping (connections, overlays, hover states)

**Done when:** Someone unfamiliar with the project can click through the prototype and understand what's happening without explanation.

---

### Task P2.2: Representation Mode Comparison (~2 hrs)

**What to do:**

- Take the same example data and create the trajectory view in 2 different visual styles:
  1. The "structured" mode (scatter plot or layer stack — more analytical)
  2. The "organic" mode (flowing, soft, spatial — more exploratory)
- Present them side by side with a sentence explaining the vibe of each

**Done when:** 2 versions of the same screen, same data, different visual language.

---

### Task P2.3: Visual Spec Document (~2 hrs)

## What to do:

- Create a Figma page called "Specs" with:
  - Color palette (token colors, background, text, interactive states)
  - Typography (font, sizes for labels, headings, data)
  - Component inventory: token chip, layer marker, trajectory line, attention cell, prediction row
  - Spacing/layout grid

**Done when:** A developer could look at this page and know exactly what CSS to write.

---

## Summary: Two-Week Deliverables

Who	Week 1 Deliverable	Week 2 Deliverable
Steven	Tokenization function, PCA reduction + matplotlib plot	/trace endpoint, /predict endpoint (logit lens)
Mario	Embedding extraction function, Flask skeleton + /tokenize	/attention endpoint, example JSON data, README
Both	Tutorial walkthrough notebook, environment setup	API documentation
Polly	5 wireframe screens, 3 representation sketches, research doc	Clickable prototype, 2 visual modes, spec document

## Time Budget

### Steven (16 hrs total)

Task	Hours	Week
1.1 Environment setup	1	1
1.2 Tutorial walkthrough	1.5	1
1.3 Tokenization function	1.5	1
1.5 Dimensionality reduction	2.5	1
Buffer / debugging	1.5	1
2.1 /trace endpoint	3	2
2.3 /predict endpoint	3	2

Task	Hours	Week
2.5 Documentation (shared)	1	2
Buffer / debugging	1.5	2

## Mario (16 hrs total)

Task	Hours	Week
1.1 Environment setup	1	1
1.2 Tutorial walkthrough	1.5	1
1.4 Embedding extraction	2	1
1.6 Flask skeleton + /tokenize	2	1
Buffer / debugging	1.5	1
2.2 /attention endpoint	3	2
2.4 Example data generation	2	2
2.5 Documentation (shared)	1	2
Buffer / debugging	2	2

## Polly (16 hrs total)

Task	Hours	Week
P1.1 Research existing projects	2	1
P1.2 User flow wireframes	3	1
P1.3 Representation sketches	3	1
P2.1 Clickable prototype	4	2
P2.2 Representation comparison	2	2
P2.3 Visual spec document	2	2

---

## Blocker Protocol

 If you're stuck, message Gabrielle.

Seek help when debugging. Common blockers:

- **TransformerLens won't install:** Try `pip install git+https://github.com/TransformerLensOrg/TransformerLens.git`
- **Model download fails:** Check internet connection, try a different network
- **Shape mismatch errors:** Print `.shape` of every tensor. GPT-2 small dimensions:  
vocab=50257, d\_model=768, n\_layers=12, n\_heads=12
- **Flask CORS errors:** Make sure `flask-cors` is installed and `CORS(app)` is in `app.py`

---

*Created: 2026-02-15*

*Sprint: Feb 17 – Feb 28, 2026*